

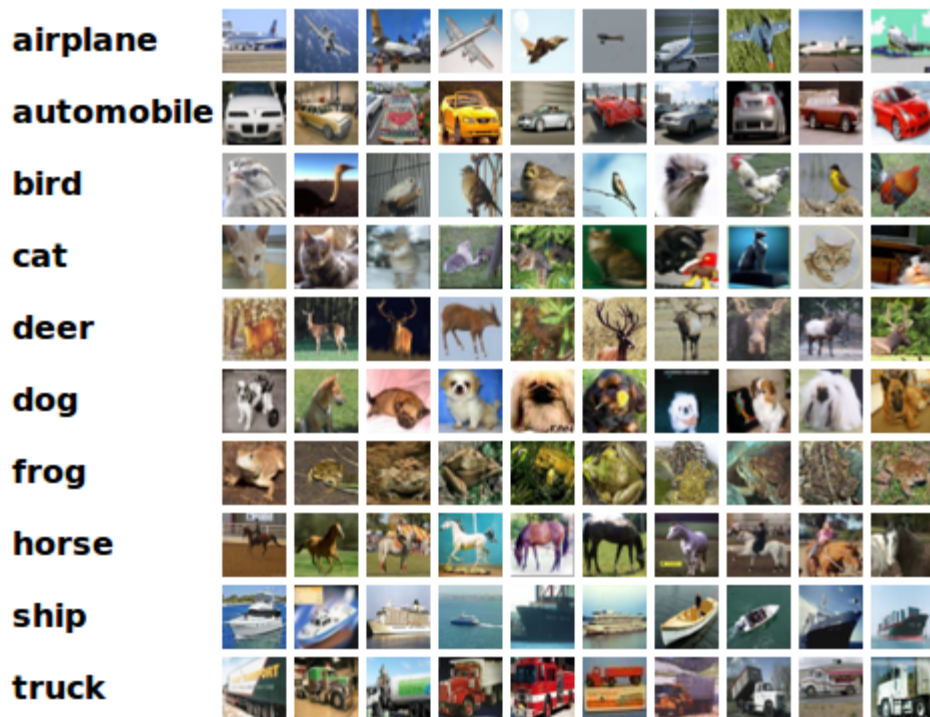
The American University in Cairo

Department of Computer Science and Engineering

CSCE 4930 – Practical Machine Deep Learning

Dr. Mohamed Moustafa	Assignment 2 [10%]	Fall 2017
----------------------	--------------------	-----------

Implement a multi-layer fully connected Neural Network (NN) classifier that **best** recognizes the 10 different classes in the CIFAR-10 dataset.



Details:

- Download the dataset from <http://www.cs.toronto.edu/~kriz/cifar.html>
The CIFAR-10 dataset consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. There are 50000 training images and 10000 test images.
- Implement YOUR OWN NN such that you can dynamically change the architecture, e.g., number of hidden layers, number of neurons per layer, activation functions, etc...
- Implement YOUR OWN gradient descent learning algorithm (preferably using the computation graph method)

- d) Use a validation set (or n-crossfold) to fine tune your hyper-parameters.
- e) Use the 10000 testing set to report the accuracy of your classifier .

You are expected to deliver: (one file: first.lastname.assignment2.zip)

1. source code of your implementation. **[3%]**
2. short report containing:
 - a) Explanation of any data-preprocessing you have used.
 - b) Explanation of how did you choose your network architecture and how did you fine tune your hyper-parameters, e.g., learning rate, regularization, etc...(explanation associated with some numerical evidence is preferred) **[1%]**
 - c) a plot of your best training and validation losses (Y axis) versus number of epochs (X axis) showing when did you stop training. **[2%]**
 - d) Correct Classification Rate of each of the 10 classes separately. CCR_n is the ratio of the number of correctly classified images in class n divided by the total number of test images in class n (which is 1000). Report CCR_n for NN versus LLS classifier (from assignment 1). **[1%]**
 - e) Average Correct Classification Rate (ACCR) using the 10000 testing set. ACCR is the number of correctly classified images divided by the total number of testing images (which is 10000). **[max(0,30(ACCR-0.5))%]**

- Human ACCR for this dataset is known to be ~ 0.94 ¹

¹<http://karpathy.github.io/2011/04/27/manually-classifying-cifar10/>